

PROJECT BACKGROUND

As the Database Administrator for a SaaS environment managing over 1,000 tenant databases, I proposed and led a project to modernize our aging SQL Server infrastructure. The existing architecture included several legacy servers, with a mix of tenants, some of which were contractually required to remain isolated on dedicated physical servers, and others were spread across three overburdened, aging machines.

I developed and pitched a plan to consolidate the shared servers into a single, modernized system, with key objectives in mind:

- Reducing maintenance and physical footprint
- Lowering power consumption
- Dramatically improving performance and reliability

For the specs-inclined:

Legacy servers:

- Processor – Quad-core Xeon
- OS Storage – 2x 7.2k HD in RAID1
- Data Storage – 4x 7.2k HD in RAID5
- RAM – 32GB DDR3

New consolidated server:

- Processors – Dual 8-core Xeons
- OS Storage – 2x SSD in RAID1
- Data Storage – 4x SSD in RAID10
- RAM – 256GB DDR4

The upgraded system would host approximately half of our tenant databases, requiring a carefully orchestrated, large-scale migration.

To streamline not only this migration but future moves as well, I proactively developed a reusable stored procedure to handle tenant database migrations on-demand. This tool automated the end-to-end process with pure T-SQL and minimal downtime.

While I do have the complete procedure code, I cannot share it due to ethical, legal, and operational security concerns. Also, the script spans over 350 lines and contains account credentials, making it impractical to present in a public-facing context. However, I will provide a full outline of the process as designed and implemented, without the 8 extra pages.

DESIGN CONSIDERATIONS

The SaaS environment consists of tenant containers distributed across multiple servers, along with a container and user registration database hosted on a dedicated server. The registration table includes a pointer indicating which server currently hosts each tenant's database.

The stored procedure was designed with the following input parameters:

- Source server
- Target server
- Backup storage network path (NAS device)
- SQL filter to determine which databases to include
- Flag indicating whether to retain or drop the source database upon success

PROCESS

1. Create Linked Servers – Establish linked servers with sa-level access for both the source and target servers.
2. Map Network Location on Source and Target Servers – Use xp_cmdshell to map the shared network location (NAS).
 - 2.1 If xp_cmdshell fails to execute, log the error and exit the process.
 - 2.2 If mapping fails, attempt to delete and recreate the network location.
 - 2.3 If the second attempt fails, log the error and exit the process.
3. Retrieve Default Storage Paths – Query the target server to determine default data and log file locations.
4. Build Migration Dataset – Generate table variable containing the list of databases, based on the user-defined query filter.
5. Migrate Databases – Iterate through the list with a cursor and perform the following steps for each:
 - 5.1 Retrieve database metadata (names, file paths) from source server.
 - 5.2 Construct dynamic SQL statements for subsequent operations.
 - 5.3 Check if the database already exists on the target server.
 - If it exists, log the conflict and continue to the next database.

5.4 Confirm the database exists on the source server.

- If it does not exist, log the conflict and continue to the next database.

5.5 Disable proprietary rule and notification processing agent, restrict user access.

5.6 Perform full backup to NAS using xp_cmdshell and SQLCMD.

5.7 Restore the database on the target server using the same method.

5.8 Connect the application service account with db_owner role.

5.9 Validate connectivity and permissions using the service account.

- If access fails, log the issue, re-enable the original source container, continue to next database.

5.10 Update the container registration table with the new server pointer.

5.11 Drop the database from the source server if enabled in the execution call.

5.12 Log success, then proceed to the next database in the cursor.

6. Cleanup – Remove temporary objects.

6.1 Remove mapped network drives from both servers.

6.2 Drop both linked servers.

USE CASE – OVER 500 DATABASES MIGRATED IN LESS THAN AN HOUR

The initial run of the migration process was completed successfully, with no unforeseen issues. Over 500 tenant databases were moved to the new server in under an hour. Thanks to the compartmentalized approach—disabling, migrating, and re-enabling one database at a time—average service interruption per tenant was less than 10 seconds.

Customers immediately reported noticeable improvements in application responsiveness. Performance monitoring confirmed those observations, revealing more than a 400% improvement in overall system performance. Despite the new server hosting three times as many databases as the individual legacy servers, backup jobs completed in nearly the same timeframe, limited only by NAS and network throughput.

TAKEAWAYS

This project demonstrates the value of combining foresight, technical fluency, and process optimization to solve complex infrastructure challenges with minimal disruption to the business. The migration of over 500 tenant databases—executed in under an hour with near-zero downtime—was not just a technical win but a strategic one. It aligned infrastructure performance with business growth, reduced hardware sprawl, and immediately improved customer satisfaction.

From a design perspective, the reusable stored procedure was built not just for this specific project but with future migrations in mind. It abstracts the complexity of linked servers, session handling, service account provisioning, and dynamic SQL execution into a repeatable, auditable process. The flexibility of the parameters allowed for on-demand execution as business needs evolved—whether it be another full server migration, or just a single database.

The experience also reinforced the importance of designing for resilience. By building in layered error handling, conditional logic, and comprehensive logging, the procedure ensured that failures were both rare and recoverable. It reduced the possibility of mistakes, reduced the reliance on manual intervention, and improved confidence in our ability to scale.

This effort reflects one of my core professional values: building tools and processes that outlive the immediate use case, whether developing migration utilities, ETL pipelines, or reporting systems.

Ultimately, this project is a great example of the kind of work I enjoy most—deep in the guts of SQL Server, solving the kind of problems people only realize they have when things go sideways. It is where I thrive: managing complexity, removing fragility, and making data infrastructure an asset instead of a bottleneck.